



Vlaanderen
is supercomputing

New Open Ondemand Features

Wouter Lampaert
Ehsan Moravveji

VSC User Day - 13 June 2024

Overview

OnDemand Upstream

- OnDemand at KU Leuven
- Globus integration

Updated Apps

- RStudio Server
- JupyterLab
- NVidia RAPIDS

New Apps

- VNC-based apps
 - Matlab
 - ParaView

Beyond OOD

- Scaling up

Overview

OnDemand Upstream

- OnDemand at KU Leuven
- Globus integration

Updated Apps

- RStudio Server
- JupyterLab
- NVidia RAPIDS

New Apps

- VNC-based apps
 - Matlab
 - ParaView

Beyond OOD

- Scaling up

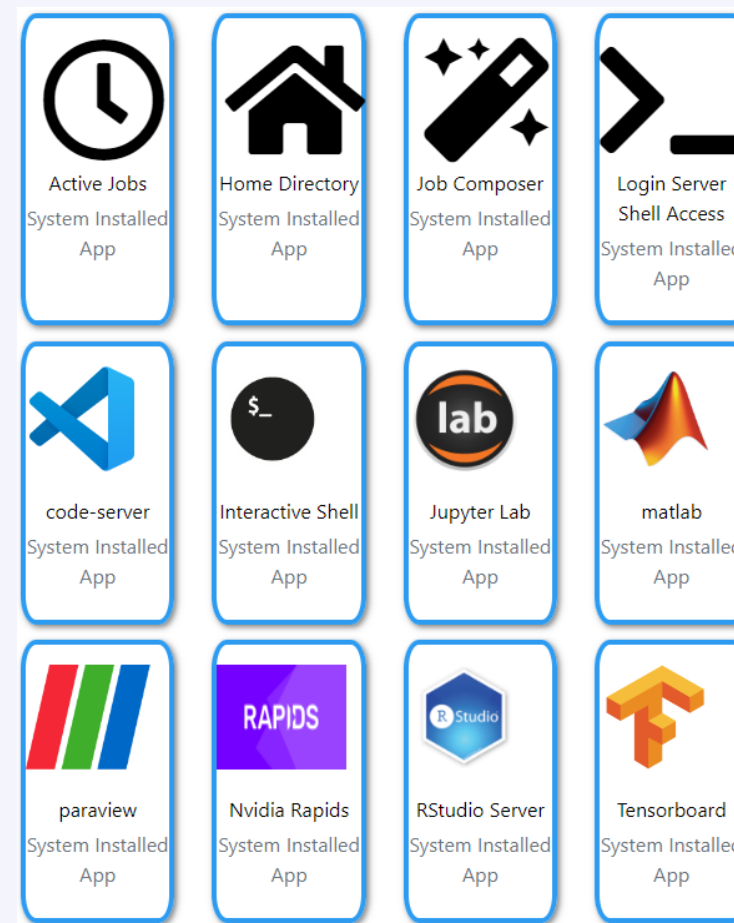
What is Open OnDemand?

Upstream

- Web-based Access to HPC facilities
More info: [OOD website](#), [OOD documentation](#)
- Latest release: 3.1.1

OOD@KULEuven

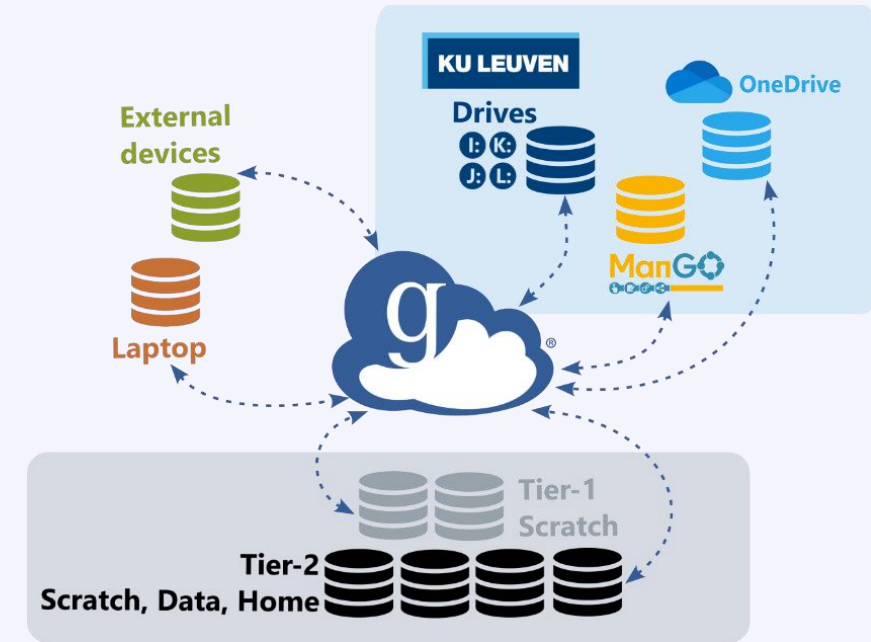
- Login via <https://ondemand.hpc.kuleuven.be>
- Authentication: MFA (no SSH keys needed!)
- Basics:
File/Folder browser, File editor, File transfer (see later)
- Native Apps: Active Jobs, Job Composer, Login & Interactive Shell
- Extra Apps (running as jobs):
Code Server, RStudio, JupyterLab, RAPIDS, Tensorboard, MATLAB, ParaView
- Last release: Yesterday!




Globus

About

- Interface to easily transfer (large) files and directories between:
 - VSC clusters
 - Remote storage (e.g. K- and L-drive)
 - Tier-1 data/ManGO
 - Remote servers
 - Personal computer
- Stable and fast transfers with restart options
- Accessible through globus.org or via the Ondemand File Browser
- Extensive user documentation available on our [Globus VSC-Docs pages!](#)
- Do you still need FileZilla/WinSCP?



Globus: using it from Open Ondemand



How-to

- Go to 'Files'
- Choose 'Data Directory' (or any sub-directory)
- Click on the 'Globus' button
- Log in to Globus with your KU Leuven ID (if asked)
- Sub-directory will be the starting location

Log in to use Globus Web App

Use your existing organizational login

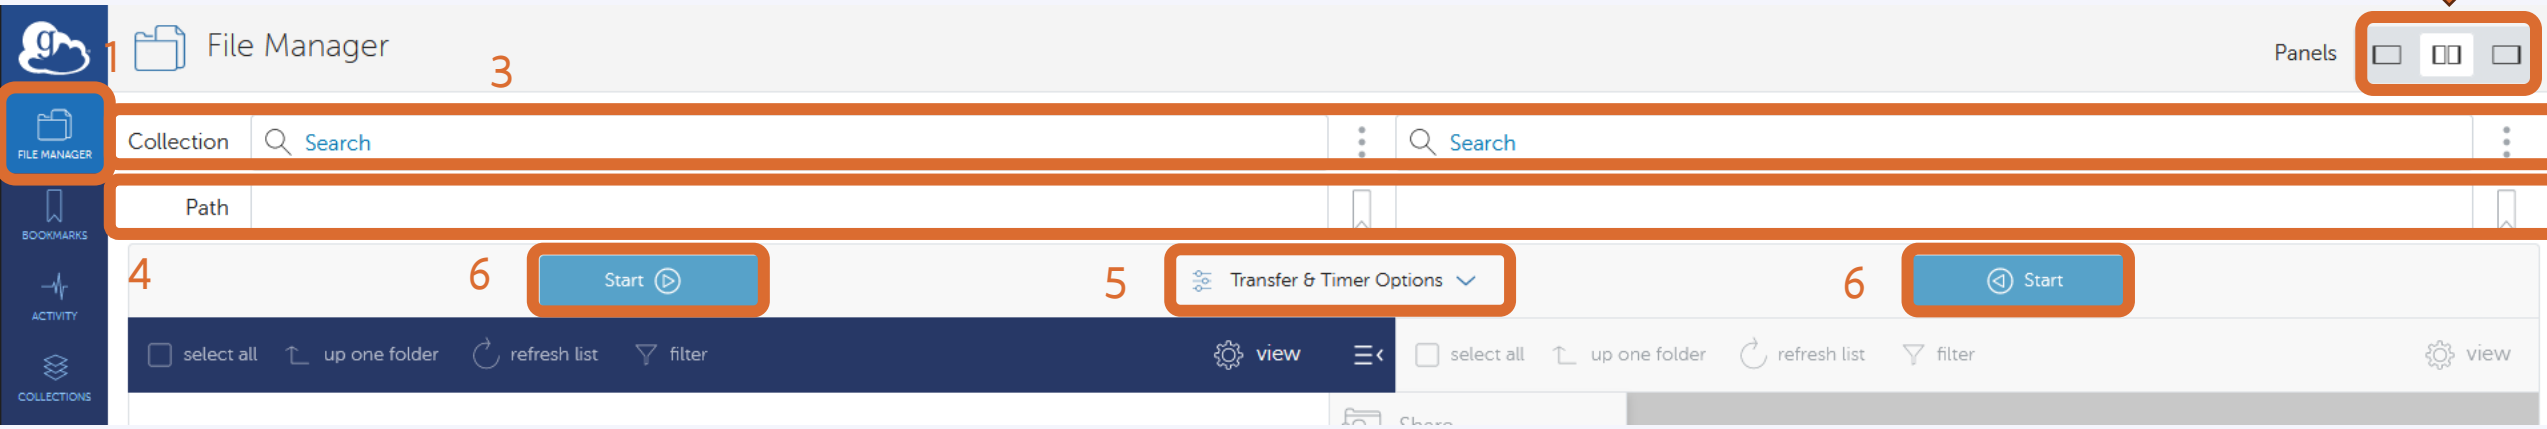
e.g., university, national lab, facility, project

KU Leuven / KU Leuven Association

By selecting Continue, you agree to Globus [terms of service](#) and [privacy policy](#).

Continue

Globus: how to use it



1. Select 'File Manager'
2. Choose the two-pane view
3. Select two collections to transfer between.
 - Have a look at our list of managed collections on our [Globus VSC-Docs pages](#)
4. Choose the two directories to transfer between
5. Set your transfer and timer options
 - Select syncing if you do not want to overwrite existing files
 - Select quality check protocol
 - Set name for transfer
6. Start the transfer

Steps

Globus: local endpoints

About

- Easily transfer to/from a machine without a managed endpoint
 - Personal computer
 - Local server in department
 - Other HPC clusters without Globus endpoints
 - ...
- While the machine/connection to a remote machine is live, the endpoint is available
- Install [Globus Connect Personal](#) to use on any machine

Overview

OnDemand Upstream

- OnDemand at KU Leuven
- Globus integration

Updated Apps

- RStudio Server
- JupyterLab
- NVidia RAPIDS

New Apps

- VNC-based apps
 - Matlab
 - ParaView

Beyond OOD

- Scaling up

RStudio Server

- RStudio installed on the OS
- Latest R was automatically pulled (now v. 4.4.0)
- But: no extra R packages
- User package management: difficult and fragile
- Possible inconsistencies w.r.t. different archs

Now

- RStudio as a module
- Using R/4.2.2-foss-2022b module
- Extra (>1,100) packages from CRAN repo
- Extra (>400) packages from Bioconductor repo
- Consistent package installation w.r.t. different archs

Soon

- Latest RStudio with R/4.3.3-gfbbf-2023b
- Latest R bundle from CRAN Bioconductor

Extra

Toolchain year and R version

2022b: Using R/4.2.2-foss-2022b

- ☒ Load R-bundle-CRAN module for additional packages
- ☒ Load R-bundle-Bioconductor module for additional packages

```
$ module load R-bundle-CRAN/2023.12-foss-2022b
$ module load R-bundle-Bioconductor/3.16-foss-2022b-R-4.2.2
$ R
> library(DESeq2)
...
> library(Seurat)
Attaching SeuratObject

Attaching package: 'Seurat'
...|
```

JupyterLab

- JupyterLab installed on the OS
- Python (pre-built) from a container
Python version: 3.9
- Too dry: no extra Python packages
- User package management: possible with pip
Fiddling with “–prefix” and “Pre-run scriptlet” needed

Before

- User packages are architecture-dependent

- JupyterLab as a module
- Choose from 2021a, 2022a, 2023a toolchains
- Extra SciPy-bundle and Matplotlib modules
- Consistent package installation w.r.t. different archs

Now

- Python v3.11.5 from 2023b toolchain
- PyTorch and TensorFlow modules
- Resource Utilization Dashboard

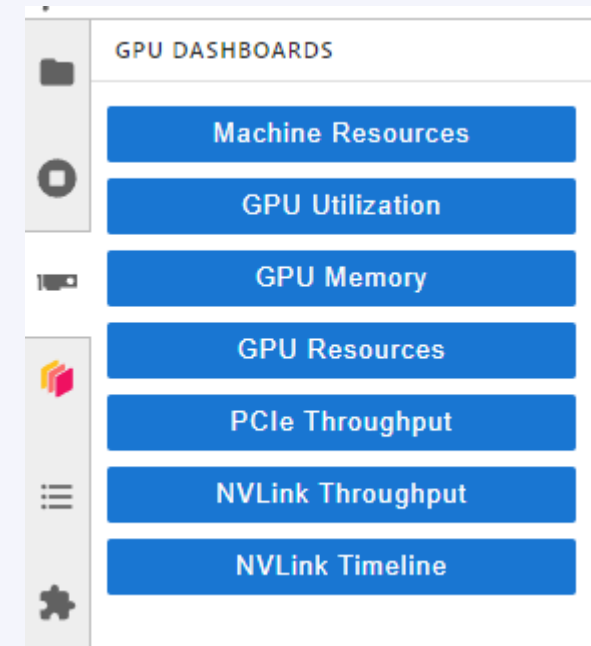
Soon

Toolchain year

2023a

☐ Load the SciPy-bundle module with FOSS

☐ Load the matplotlib module with FOSS



Nvidia RAPIDS

About

- See <https://rapids.ai/>
- GPU-accelerated data structures and algorithms
- Drop-in replacement for e.g. Pandas, NetworkX, scikit-learn, scikit-image, and more
- For Data Science, Machine Learning and more

Setup

- RAPIDS launched from container
- Converted from Docker to Apptainer
- Choosing RAPIDS version
Container names follow RAPIDSAI convention
- Stable releases are included
- Now: Using JupyterLab
- Kernel: Python 3 (ipykernel)

FYI

- RAPIDS v.24.04 is incompatible with P100 GPUs

RAPIDS module

RAPIDS/24.04-cuda12.2-py3.11

RAPIDS/22.10-CUDA-11.5-runtime-Ubuntu-18.04-Python-3.9-singularity
RAPIDS/23.06-cuda11.8-runtime-ubuntu22.04-py3.10-apptainer-jupyter-1.24.0
RAPIDS/24.04-cuda12.2-py3.11

Example

```
import numpy
import pandas
import cudf

# instantiate data directly in GPU memory
cdf = cudf.DataFrame({'indx': numpy.arange(10),
                      'vals': numpy.random.random(10)})

arr = numpy.ndarray((101, 101, 21), ...) # in CPU memory
garr = cudf.DataFrame.from_records(arr, ...) # in GPU memory

pdf = pandas.DataFrame(...) # native pandas
gdf = cudf.from_pandas(pdf) # convert from Pandas
```

Nvidia RAPIDS

Benchmark

- [cudf benchmark](#) on github
- GPU: 1x Nvidia V100
CUDA v. 12.0
- CPU: 4x CascadeLake core
Frequency: 2.6GHz
Memory: 20GB/core
- RAPIDS v. 24.04

Metrics

- Measuring runtime
- Only for function calls
- Data movement (copy, allocation) on GPU is not measured
- Unfair comparison

○ $\text{Runtime_CPU} / \text{Runtime_GPU}$

Speedup

Nvidia RAPIDS

Numerics

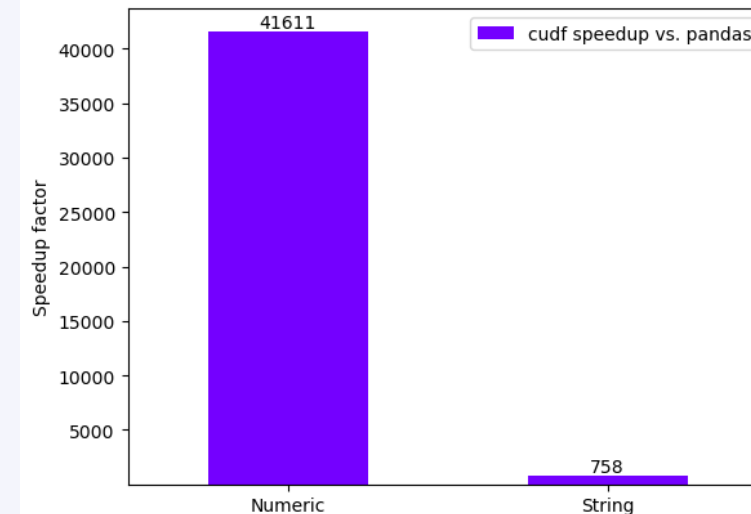
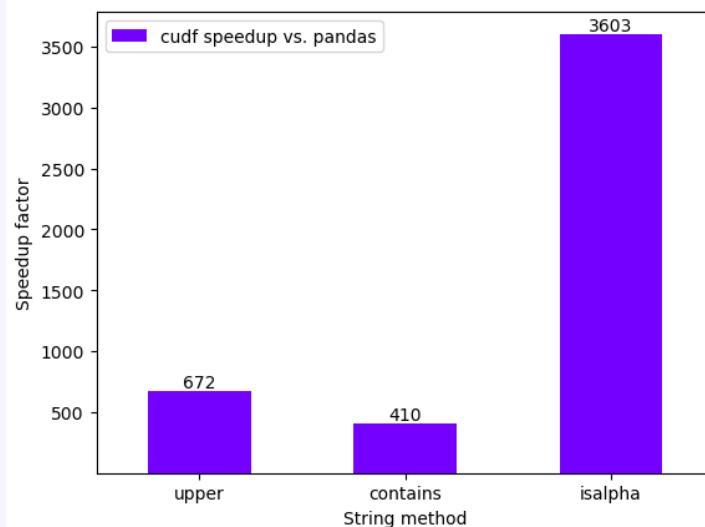
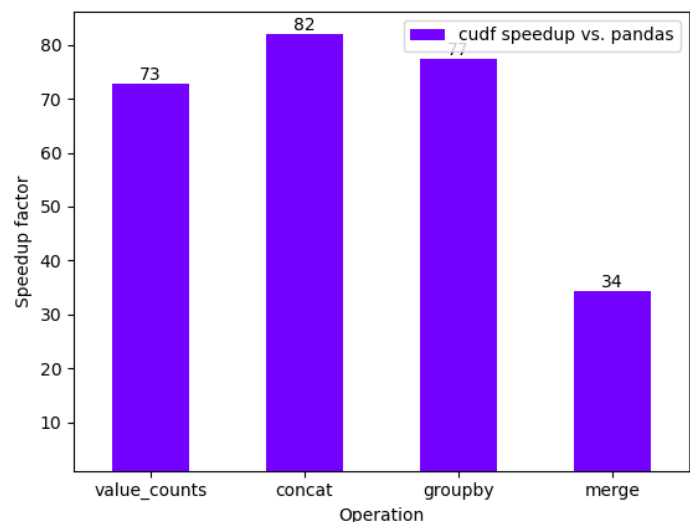
- Counting
- Concatenating
- Group-by comparison
- Merge datasets

Strings

- Uppercase
- Contains
- Isalpha

UDF

- UDF: User-defined Function
- Numerics and Strings
- Without JIT overhead
(=after warmup)



Overview

OnDemand Upstream

- OnDemand at KU Leuven
- Globus integration

Updated Apps

- RStudio Server
- JupyterLab
- NVidia RAPIDS

New Apps

- VNC-based apps
 - Matlab
 - ParaView

Beyond OOD

- Scaling up

VNC-based apps

About

- Running interactive apps on a remote desktop in the background
- (Faster) integration of interactive apps into Open Ondemand
 - Taking user requests, but... this will need review from our side
- VNC-hosted apps are minimalistic and contain a single app
 - Easier maintenance of each app
 - Problem with one app would not block access to other apps
- New apps:
 - MATLAB
 - ParaView



VNC-based apps: new options

Resolution

- Available on the form
- Pre-defined options (Full HD, 2K, 4K)
- Fine-tune with the sliders at startup, but less recommended

Sharing

- Option to share your session with a link (view-only)

matlab (56265758) 1 node | 4 cores | Running

Host: [_r25il3n17.genius.hpc.kuleuven.be](#) Cancel

Created at: 2024-06-07 14:00:48 CEST

Time Remaining: 59 minutes

Session ID: [7edccf33-e769-417a-97b3-481ac4f197ff](#)

Compression 0 (low) to 9 (high)

Image Quality 0 (low) to 9 (high)

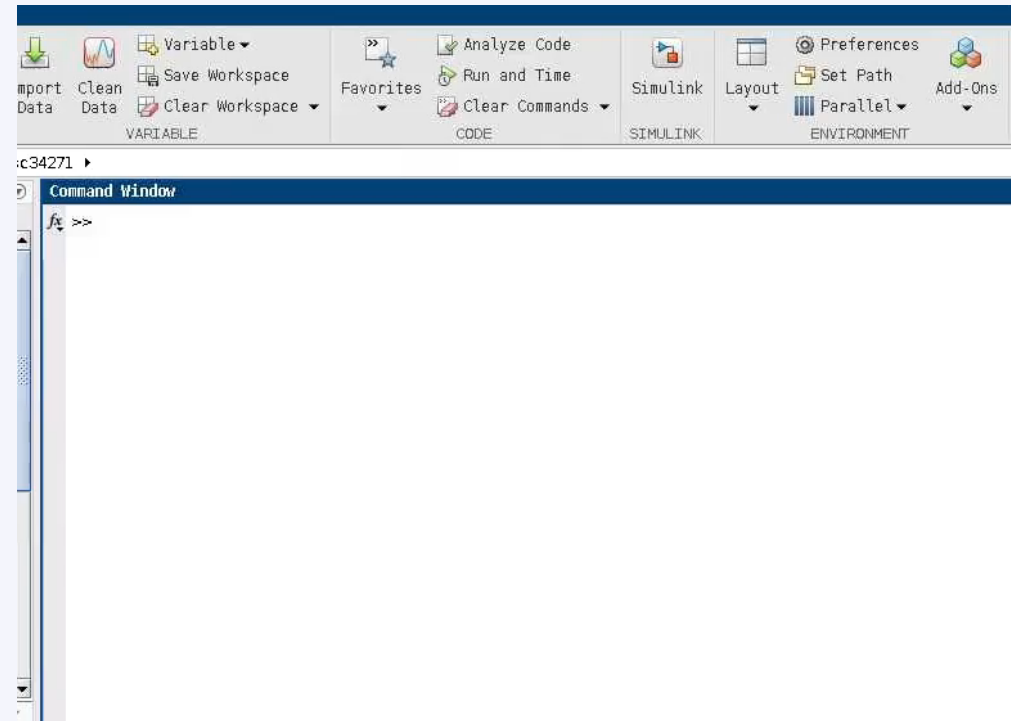
Launch matlab

View Only (Share-able Link)

VNC-based apps: MATLAB

About


- Current supported versions:
 - MATLAB/2022a
 - MATLAB/2022b
 - Available on Genius and wICE
 - Access: join `lli_matlab` group
- Allows running the GUI on compute nodes
- Improved performance compared to NX session (login node)
- GPU support



VNC-based apps: ParaView

About

- Current supported versions:
 - 5.6.1
 - 5.9.1-FOSS
 - 5.9.1-Intel
 - 5.10.1-FOSS
 - On both Genius and WICE



Overview

OnDemand Upstream

- OnDemand at KU Leuven
- Globus integration

Updated Apps

- RStudio Server
- JupyterLab
- NVidia RAPIDS

New Apps

- VNC-based apps
 - Matlab
 - ParaView

Beyond OOD

- Scaling up

Growing out of OOD

IF

- ✓ your dataset is ready
- ✓ your workflow is well tested
- ✓ your scripts are mature
- ✓ Longer walltime needed

THEN

Go to “batch” mode

- Your input data stays the same
- Your scripts stay the same

Python

```
#!/bin/bash -l
#SBATCH --account=lp_science
#SBATCH --cluster=wice
#SBATCH --nodes=1 --ntasks=1
#SBATCH --time=2-00:00:00

module load JupyterLab/4.0.5-GCCcore-12.3.0

jupyter execute prepare.ipynb analyse.ipynb postproc.ipynb
```

Load extra
modules

Run multiple
notebooks

Growing out of OOD

IF

- ✓ your dataset is ready
- ✓ your workflow is well tested
- ✓ your scripts are mature
- ✓ Longer walltime needed

THEN

Go to “batch” mode

- Your input data stays the same
- Your scripts stay the same

R

```
#!/bin/bash -l
#SBATCH --account=lp_science
#SBATCH --cluster=wice
#SBATCH --nodes=1 --ntasks=1
#SBATCH --time=2-00:00:00

module load R/4.2.2-foss-2022b
module load R-bundle-CRAN/2023.12-foss-2022b
Rscript clustering.R
```

Growing out of OOD

IF

- ✓ your dataset is ready
- ✓ your workflow is well tested
- ✓ your scripts are mature
- ✓ Longer walltime needed

THEN

Go to “batch” mode

- Your input data stays the same
- Your scripts stay the same

RAPIDS

```
#!/bin/bash -l
#SBATCH --account=lp_science
#SBATCH --cluster=wice --partition=gpu_a100
#SBATCH --nodes=1 --ntasks=1 --gpus-per-node=1
#SBATCH --time=2-00:00:00

module load RAPIDS/24.04-cuda12.2-py3.11

apptainer exec --nv $RAPIDSIMAGE python3 data_science.py
```

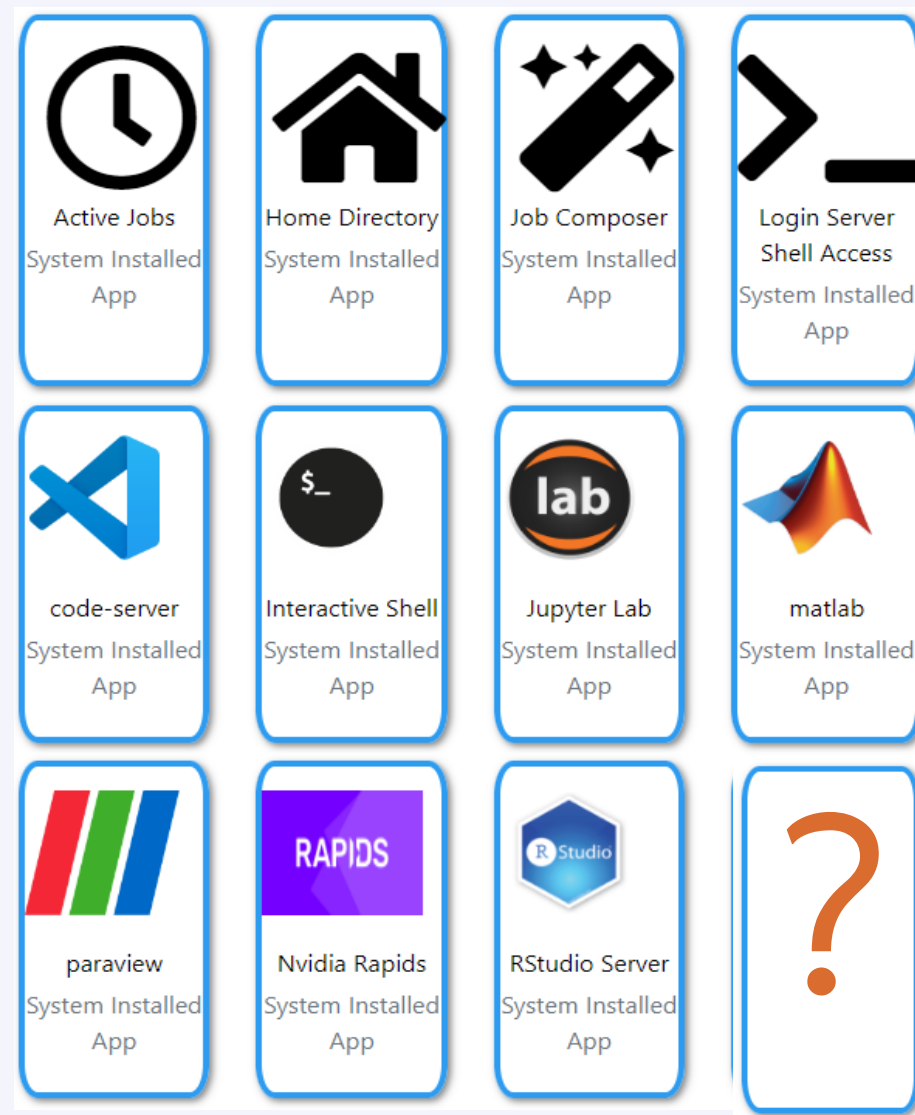
Contains RAPIDS
conda env.

Python script
must be
executable

Finally

Work in progress

- You need a new app or feature?
- Contact us via hpcinfo@kuleuven.be
- Development takes time
testing -> quality -> production
- Next candidates:
Ansys FLUENT
Stata
BAND



*Stay Connected
to VSC*

Linked  [®]

